

Exercise: Two-way ANOVA in R

Stephan Huber

2024-06-04

This is an exercise where data management with the `dplyr` functions `pivot_longer`, `rename`, and `bind_rows` is practiced. Moreover, I exemplify how an ANOVA analysis can be executed with R. Doing so I refer to the content of Childs et al. (2021, Chapter 27). All data to this exercise can be found in Huber (2024).

Table of contents

1	Read	2
2	Set up the R session	2
3	Create a data frame	2
4	Make the data “long”	3
5	Descriptive statistics	4
6	ANOVA	5
7	Diagnostics	6
8	Interaction diagram	7
9	Multiple comparison tests	11
10	Drawing conclusions and presenting results	12
	References	15

1 Read

Read Childs et al. (2021): [27.2 Competition between Calluna and Festuca](#).

Our goal is to learn how to work with two-way ANOVA models in R, using an example from a plant competition experiment. The work flow is very similar to one-way ANOVA in R. We'll start with the problem and the data, and then work through model fitting, evaluating assumptions, significance testing, and finally, presenting the results.

2 Set up the R session

Download and open the script that you find [here](#). This script contains all the code that is shown below. The solutions to the exercises can be found [here](#).

```
rm(list = ls())

if (!require(pacman)) install.packages("pacman")
pacman::p_load(tidyverse, rstatix, ggpubr, agricolae)
```

3 Create a data frame

Plants have an optimal soil pH for growth, and this varies between species. Consequently we would expect that if we grow two plants in competition with each other at different pH values the effect of competition might vary according to the soil pH. In a recent study the growth of the grass *Festuca ovina* (Sheep's Fescue) in competition with the heather *Calluna vulgaris* (Ling) was investigated in soils with different pH. *Calluna* is well adapted to grow on very acidic soils such as on the Millstone grit and blanket bogs around Sheffield. *Festuca* grows on soils with a much wider range of pH. We might hypothesise that *Calluna* will be a better competitor of *Festuca* in very acid soils than in moderately acid soils. Here are the data: The column (Weight) contains the *Festuca* dry weights, (pH) contains the codes for the pH treatment (levels: pH3.5, pH5.5), the column (Calluna) contains the codes for the presence or absence of *Calluna* (levels: Present, Absent).

```
data_present <- data.frame(
  Condition = rep(c("Calluna Present"), each = 5),
  `pH 3.5` = c(2.76, 2.39, 3.54, 3.71, 2.49),
  `pH 5.5` = c(3.21, 4.10, 3.04, 4.13, 5.21),
  check.names = FALSE
)
data_present
```

	Condition	pH 3.5	pH 5.5
1	Calluna Present	2.76	3.21
2	Calluna Present	2.39	4.10
3	Calluna Present	3.54	3.04
4	Calluna Present	3.71	4.13
5	Calluna Present	2.49	5.21

```
data_absent <- data.frame(
  Condition = rep(c("Calluna Absent"), each = 5),
  `pH 3.5` = c(4.10, 2.72, 2.28, 4.43, 3.31),
  `pH 5.5` = c(5.92, 7.31, 6.10, 5.25, 7.45),
  check.names = FALSE
)
data_absent
```

	Condition	pH 3.5	pH 5.5
1	Calluna Absent	4.10	5.92
2	Calluna Absent	2.72	7.31
3	Calluna Absent	2.28	6.10
4	Calluna Absent	4.43	5.25
5	Calluna Absent	3.31	7.45

4 Make the data “long”

Read the [R documentation](#) of the function `bind_rows` which is part of the `dplyr` package.

Combine the objects `data_present` and `data_absent`.

	Condition	pH 3.5	pH 5.5
1	Calluna Present	2.76	3.21
2	Calluna Present	2.39	4.10
3	Calluna Present	3.54	3.04
4	Calluna Present	3.71	4.13
5	Calluna Present	2.49	5.21
6	Calluna Absent	4.10	5.92
7	Calluna Absent	2.72	7.31
8	Calluna Absent	2.28	6.10
9	Calluna Absent	4.43	5.25
10	Calluna Absent	3.31	7.45

Read Wickham & Grolemund (2023): [5.3 Lengthening data](#).

Rearrange the data so that they look like this:

```
festuca
```

```
# A tibble: 20 x 3
  Calluna      pH    Weight
  <chr>      <chr>   <dbl>
1 Calluna Present pH 3.5    2.76
2 Calluna Present pH 5.5    3.21
3 Calluna Present pH 3.5    2.39
4 Calluna Present pH 5.5    4.1
5 Calluna Present pH 3.5    3.54
6 Calluna Present pH 5.5    3.04
7 Calluna Present pH 3.5    3.71
8 Calluna Present pH 5.5    4.13
9 Calluna Present pH 3.5    2.49
10 Calluna Present pH 5.5    5.21
11 Calluna Absent  pH 3.5    4.1
12 Calluna Absent  pH 5.5    5.92
13 Calluna Absent  pH 3.5    2.72
14 Calluna Absent  pH 5.5    7.31
15 Calluna Absent  pH 3.5    2.28
16 Calluna Absent  pH 5.5    6.1
17 Calluna Absent  pH 3.5    4.43
18 Calluna Absent  pH 5.5    5.25
19 Calluna Absent  pH 3.5    3.31
20 Calluna Absent  pH 5.5    7.45
```


5 Descriptive statistics

Calculate the following statistics and graphs:

```
summary_stats
```

```
# A tibble: 4 x 4
  Calluna      pH    mean  var
  <chr>      <chr>   <dbl> <dbl>
1 Calluna Absent  pH 3.5   3.37 0.818
2 Calluna Absent  pH 5.5   6.41 0.893
```

```
3 Calluna Present pH 3.5 2.98 0.371
4 Calluna Present pH 5.5 3.94 0.754
```



6 ANOVA

Use this model to calculate the ANOVA: `Weight ~ pH + Calluna + pH:Calluna`

Analysis of Variance Table

Response: Weight


	Df	Sum Sq	Mean Sq	F value	Pr(>F)
pH	1	19.9800	19.9800	28.1792	7.065e-05 ***
Calluna	1	10.2102	10.2102	14.4001	0.00159 **
pH:Calluna	1	5.3976	5.3976	7.6126	0.01397 *
Residuals	16	11.3446	0.7090		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

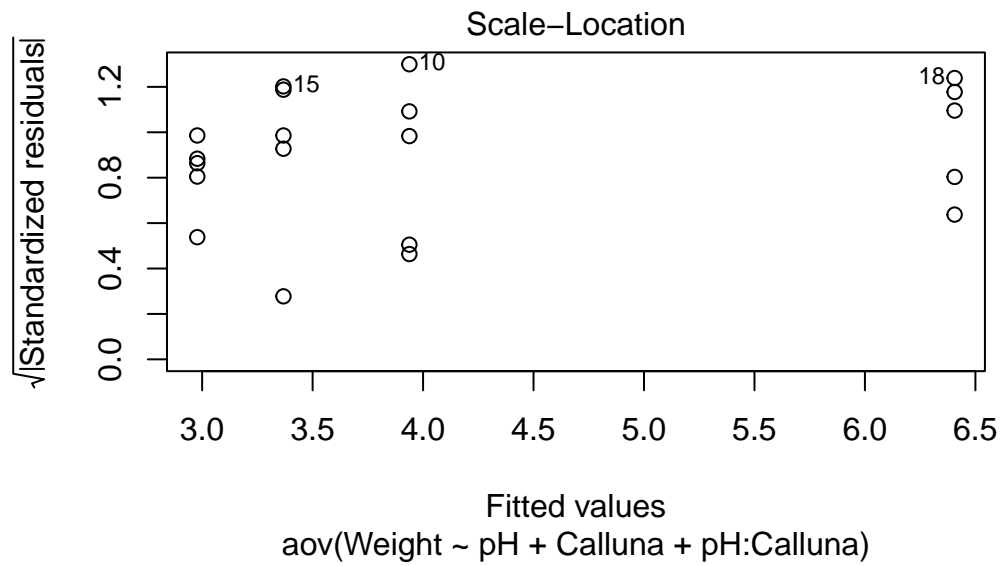
7 Diagnostics

Read Childs et al. (2021): [27.5 Diagnostics](#). Moreover, [this page](#) is worth a look. You find some alternatives R functions that you may find helpful for ANOVA diagnostics.

```
plot(festuca_model, which = 2, add.smooth = FALSE)
```

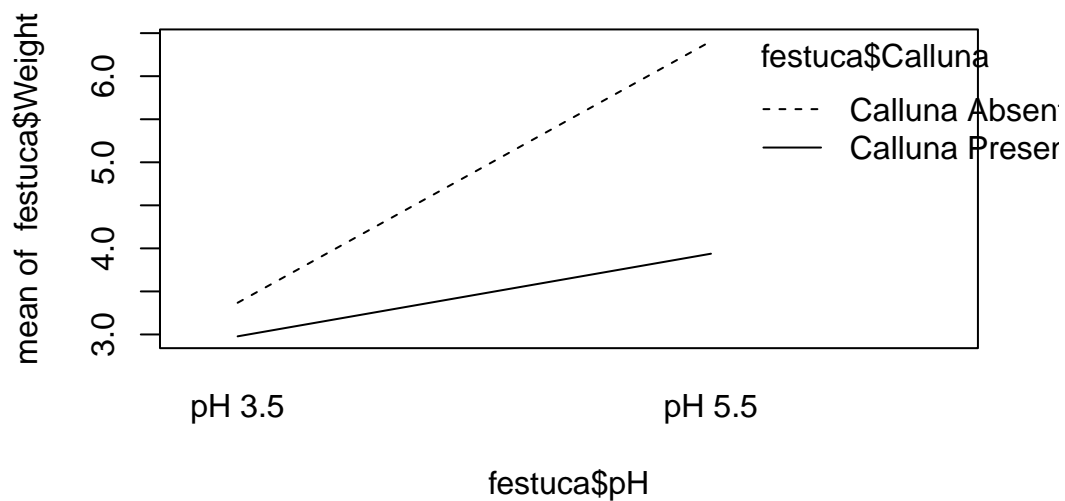


```
plot(festuca_model, which = 3, add.smooth = FALSE)
```



8 Interaction diagram

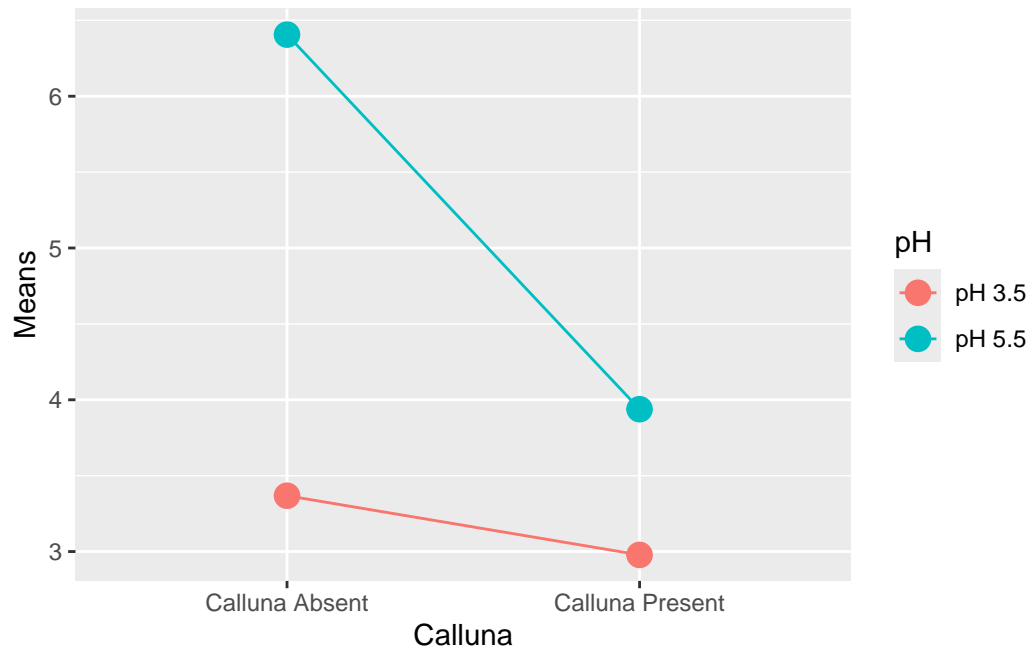
Can you use the function `interaction.plot` to create the following:



Here is a much nicer and more flexible approach to make interaction plots using **tidyverse** functions:

```
# step 1. calculate means for each treatment combination
festuca_means <-
  festuca %>%
  group_by(Calluna, pH) %>% # <- remember to group by *both* factors
  summarise(Mean = mean(Weight))
```

```
# step 2. plot these as an interaction plot
ggplot(festuca_means,
  aes(x = Calluna, y = Mean, colour = pH, group = pH)) +
  geom_point(size = 4) + geom_line()
```

Read Childs et al. (2021): [27.6.1](#) and consider the following figure:

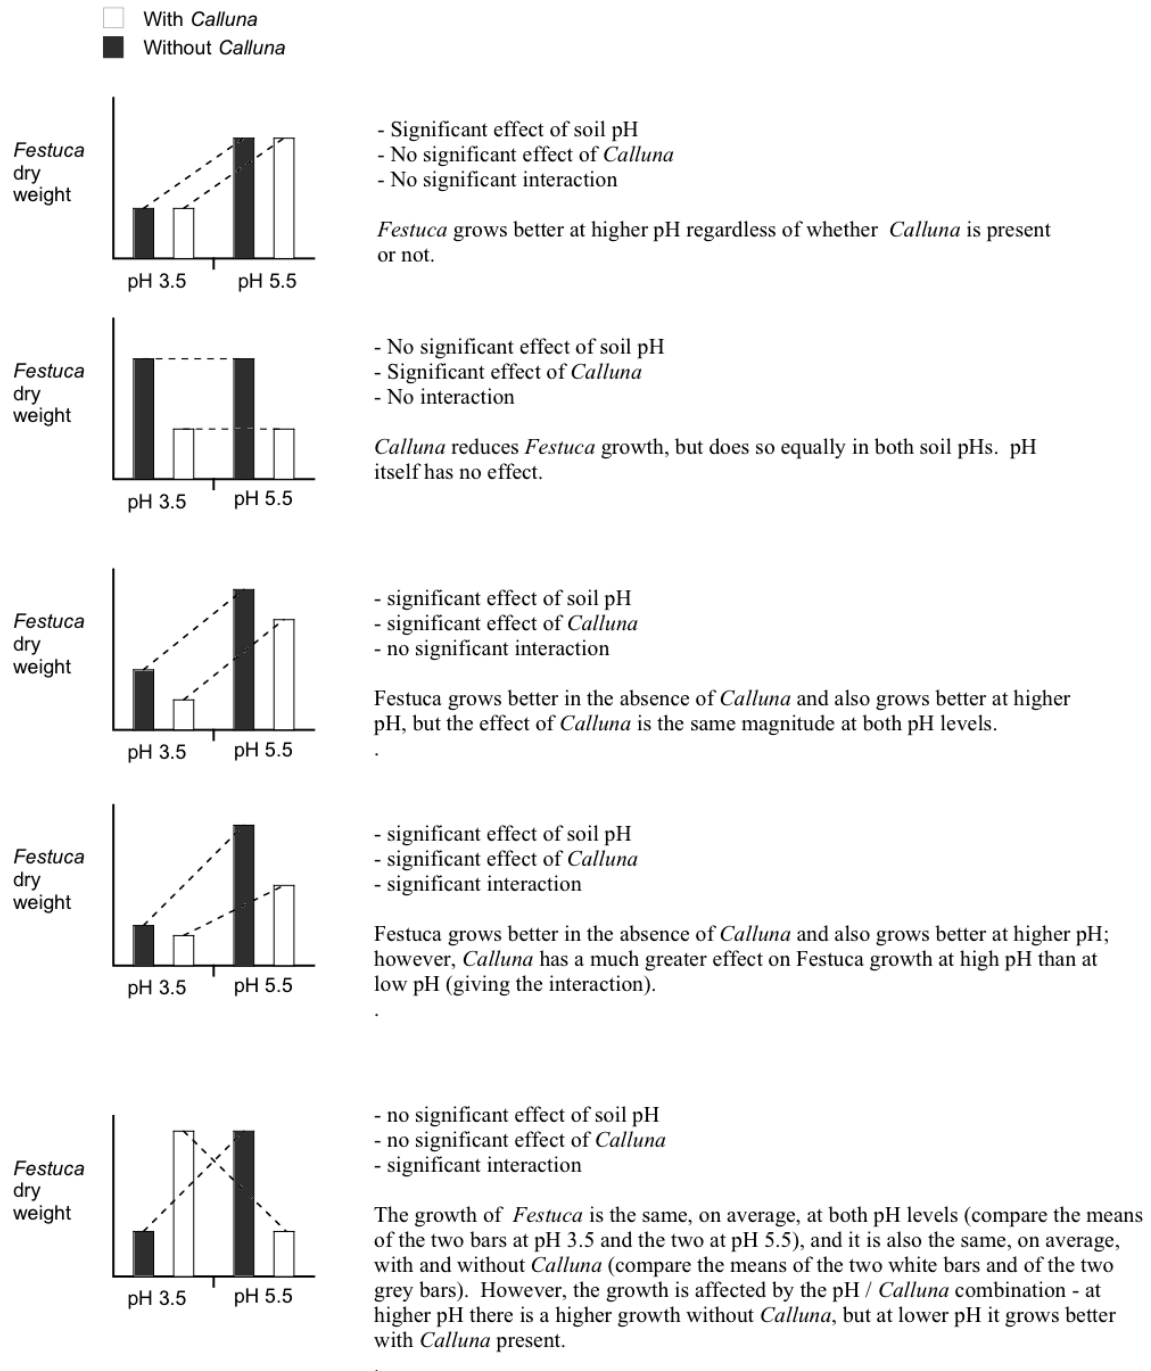


Figure 1: Understanding the model graphically

9 Multiple comparison tests

```
TukeyHSD(festuca_model, which = 'pH:Calluna')
```

Tukey multiple comparisons of means
95% family-wise confidence level

```
Fit: aov(formula = Weight ~ pH + Calluna + pH:Calluna, data = festuca)
```

```
$`pH:Calluna`
```

	diff	lwr	upr
pH 5.5:Calluna Absent-pH 3.5:Calluna Absent	3.038	1.5143518	4.5616482
pH 3.5:Calluna Present-pH 3.5:Calluna Absent	-0.390	-1.9136482	1.1336482
pH 5.5:Calluna Present-pH 3.5:Calluna Absent	0.570	-0.9536482	2.0936482
pH 3.5:Calluna Present-pH 5.5:Calluna Absent	-3.428	-4.9516482	-1.9043518
pH 5.5:Calluna Present-pH 5.5:Calluna Absent	-2.468	-3.9916482	-0.9443518
pH 5.5:Calluna Present-pH 3.5:Calluna Present	0.960	-0.5636482	2.4836482
	p adj		
pH 5.5:Calluna Absent-pH 3.5:Calluna Absent	0.0001731		
pH 3.5:Calluna Present-pH 3.5:Calluna Absent	0.8826936		
pH 5.5:Calluna Present-pH 3.5:Calluna Absent	0.7117913		
pH 3.5:Calluna Present-pH 5.5:Calluna Absent	0.0000443		
pH 5.5:Calluna Present-pH 5.5:Calluna Absent	0.0014155		
pH 5.5:Calluna Present-pH 3.5:Calluna Present	0.3079685		

```
::: {.cell}
```

```
```.r .cell-code}
```

```
HSD.test(festuca_model, trt = c("pH", "Calluna"), console = TRUE)
```

```
Study: festuca_model ~ c("pH", "Calluna")
```

HSD Test for Weight

Mean Square Error: 0.709035

pH:Calluna, means

		Weight	std	r	se	Min	Max	Q25	Q50	Q75
pH 3.5:Calluna	Absent	3.368	0.9042511	5	0.3765727	2.28	4.43	2.72	3.31	4.10
pH 3.5:Calluna	Present	2.978	0.6089089	5	0.3765727	2.39	3.71	2.49	2.76	3.54
pH 5.5:Calluna	Absent	6.406	0.9451614	5	0.3765727	5.25	7.45	5.92	6.10	7.31
pH 5.5:Calluna	Present	3.938	0.8685448	5	0.3765727	3.04	5.21	3.21	4.10	4.13

Alpha: 0.05 ; DF Error: 16

Critical Value of Studentized Range: 4.046093

Minimum Significant Difference: 1.523648

Treatments with the same letter are not significantly different.

	Weight	groups
pH 5.5:Calluna Absent	6.406	a
pH 5.5:Calluna Present	3.938	b
pH 3.5:Calluna Absent	3.368	b
pH 3.5:Calluna Present	2.978	b

...

## 10 Drawing conclusions and presenting results


Here are some code examples how the shown graphs above could be made much nicer.

```
step 1. calculate means for each treatment combination
festuca_stats <-
 festuca %>%
 group_by(Calluna, pH) %>% # <- remember to group by the two factors
 summarise(Means = mean(Weight), SEs = sd(Weight)/sqrt(n()))
```

`summarise()` has grouped output by 'Calluna'. You can override using the `.groups` argument.


```
step 1. calculate means for each treatment combination
festuca_stats <-
 festuca %>%
 group_by(Calluna, pH) %>% # <- remember to group by the two factors
 summarise(Means = mean(Weight), SEs = sd(Weight)/sqrt(n()))
```

```
step 2. plot these as an interaction plot
ggplot(festuca_stats,
 aes(x = Calluna, y = Means, colour = pH,
 ymin = Means - SEs, ymax = Means + SEs)) +
 # this adds the mean
 geom_point(size = 3) +
 # this adds the error bars
 geom_errorbar(width = 0.1) +
 # controlling the appearance
 scale_y_continuous(limits = c(2, 7)) +
 xlab("Calluna") + ylab("Festuca yield (g dry weight)") +
 # use a more professional theme
 theme_bw()
```




```
define a position adjustment
pos <- position_dodge(0.15)
make the plot
ggplot(festuca_stats,
 aes(x = Calluna, y = Means, colour = pH,
 ymin = Means - SEs, ymax = Means + SEs)) +
 # this adds the mean (shift positions with 'position =')
 geom_point(size = 3, position = pos) +
```

```
this adds the error bars (shift positions with 'position =')
geom_errorbar(width = 0.1, position = pos) +
controlling the appearance
scale_y_continuous(limits = c(2, 7)) +
xlab("Calluna") + ylab("Festuca yield (g dry weight)") +
use a more professional theme
theme_bw()
```



```
ggplot(festuca_stats,
 aes(x = Calluna, y = Means, fill = pH,
 ymin = Means - SEs, ymax = Means + SEs)) +
this adds the mean
geom_col(position = position_dodge()) +
this adds the error bars
geom_errorbar(position = position_dodge(0.9), width=.2) +
controlling the appearance
xlab("Calluna") + ylab("Festuca yield (g dry weight)")
```



## References

- Childs, D. Z., Hindle, B. J., & Warren, P. H. (2021). *APS 240: Data analysis and statistics with r*. online. <https://dzchilds.github.io/stats-for-bio>
- Huber, S. (2024). *Empirisch-wissenschaftlich arbeiten (ewa)*. GitHub repository. <https://github.com/hubchev/ewa>
- Wickham, H., & Grolemund, G. (2023). *R for data science (2e)*. Accessed January 30, 2023. <https://r4ds.hadley.nz/>